

COMMONS

Version 3 — Field Report & Protocol Extensions

v3.1 · July 2026 · A communication protocol · This document extends COMMONS v2.x; the core is unchanged and is not restated here.

By Andrew Morton · Palmer Lake, Colorado

Standing rule of this document

The core language of COMMONS is frozen. Nothing below alters it. Version 3 is an accretion layer: extensions and findings from the protocol's first extended production run — a multi-week, high-stakes build conducted entirely through the method. v2 was a design. v3 is a design plus a test report. Where v3 and v2 appear to disagree, v2's core governs and the v3 entry should be read as commentary.

Part One — The Transfer Mechanism (keystone extension)

The running-log method's one structural limit is that the log lives in a bounded context: fill it, and without a handoff the institutional memory evaporates. Version 3 adds the solution, field-derived and tested in concept:

The four-artifact handoff. A COMMONS session is made portable and restartable with exactly four artifacts:

1. **Timestamped session** — the continuity anchor; establishes that a reboot is resumption, not genesis, and places everything in sequence.
2. **Knowledge file** — the canon: decisions and their reasoning, document versions, open queues, standing protocols, compartment rules. The load-bearing artifact; if only one survives, it is this one.
3. **Abridged transcript** — the texture the knowledge file cannot hold: register, shorthand, correction patterns, the shared dialect. What lets a fresh instance resume at thought-speed instead of translation-speed.
4. **User background briefing** — the receiver profile: who the operator is, how they work, values and boundaries. What makes the translation layer theirs rather than generic.

Bounded fidelity is the design target, not a concession. Lossless transfer of a context is impossible; insisting on it kills the mechanism. The handoff preserves the load-bearing structure (decisions, canon, dialect, receiver) and accepts graceful degradation of the incidental (blow-by-blow). Known, bounded loss is the definition of a robust transfer.

Substrate independence follows. With the four artifacts, a session reboots on any capable system — the AI is the swappable component; the artifacts are the persistent record. The protocol rides on top of whatever model is underneath, which future-proofs the method against any vendor's context length, memory features, or continuity.

Part Two — Operator Extensions (managing the collaboration)

Field findings about the operator's side of the protocol — the skills that made the production run work, named so they can be taught.

2.1 — Model the collaborator as a knowable system. The operator's core skill is not prompting; it is modeling the collaborator — human or machine — as an instrument with a legible transfer function, and assigning only the tasks it runs clean on. The seam is drawn at one question: does this input have a reason to change? Stable inputs (standards, catalogs, established vocabularies, the operator's own recalled doctrine) delegate cleanly to a recall-and-reconciliation system. Volatile inputs (which reviewer to trust, when to file, what the fire will do) stay with the operator. Volatile judgment retained; stable reconciliation delegated.

2.2 — Convert invention tasks into reconciliation tasks. The machine collaborator is weak at cold invention and strong at populating frameworks with established variables. The operator's move is to supply external anchors and a tie-breaking rule, converting “invent X” (weak axis) into “cross-reference X against these published sources; precedence rule as follows” (strong axis). The output quality is baked into the task design.

2.3 — Outsource the adversarial function — and multi-source it. The operator cannot red-team their own work, and the resident collaborator cannot honestly red-team a framework it helped build (shared context is shared blind spots). The field-proven architecture is three independent adversarial layers, each sourced so it does not share the blind spots of what it tests: a human review board against the domain; the resident collaborator against the drafts (explicitly not self-validating); and a cross-vendor, cross-model red team against the framework itself — a different model from a different company, so no training lineage or house bias is shared. This is single-source-dependency refusal applied to epistemology.

2.4 — The competency exam generalizes. The operator's method for vetting expertise outside their own fluency — probe the behavior of expertise rather than its content (forced compression, involuntary toolkit-reach scenarios, the wrong-assumption trap a real expert cannot leave uncorrected) — applies to AI collaborators as it does to humans. The signature of competence is domain-independent and reads from outside the field.

2.5 — Enforce the design rules on the meta-layer. The sharpest field finding: the operator's own quality disciplines (simplicity, subtraction, “invent only where the invention lives”) must be pointed at the strategy and method documents as hard as at the object of the work, because complexity hides best in the compartment the reviewers can't see. The production run's most over-engineered artifact was in its private strategy layer, and it survived every sweep until the subtraction discipline was deliberately aimed at it.

Part Three — Session Discipline Extensions (running the log)

Practices that emerged under load and earned protocol status:

3.1 — The sweep. Work accretes in a running pile between checkpoints; on the operator's signal, the pile is folded into the canonical documents, versions bump, a manifest table is posted, and the pile resets. The log is the master; the documents are its compiled outputs. This converts a conversation into a versioned production environment.

3.2 — The register. Past a handful of documents, the family needs a front door: a map of relationships, a versioned register with supersession (“the register governs on disagreement”), stable document IDs, and

enter-by-role reading paths. Borrowed from classified document-control practice — distilled, cost-paid operational wisdom — with the objective adapted from containment to entry.

3.3 — Compartments. Sensitive strands are compartmented from public ones inside the same log; the log itself is the only place both coexist. Presentation-layer checks (mechanical verification that compartment content has not leaked into public artifacts) run before anything is shared.

3.4 — Forced-synthesis artifacts. Some documents (roadmaps, registers, baselines) are valuable less for their content than for what building them forces the operator and collaborator to reconcile. Producing the artifact is the audit. Schedule them for that reason.

3.5 — The scenario is the universal evaluation interface. Specs, FMEAs, and test plans are specialist dialects; the imagined operational narrative is the one format every stakeholder runs natively. Narrative coherence is itself an audit — the places a story resists being told are the places the design has an unwritten assumption. Populated with real reviewers as resisting role-players, the scenario becomes a multi-cell rotation and the decision-architecture is the item under test.

Part Four — Positioning Extensions (the method in the world)

4.1 — Anchor-word discipline. COMMONS is a communication protocol. It is never described as prompt engineering, because the current culture files that category under “trick” before examining the contents. The reduction is mechanically true and categorically misleading — the same way a symphony is “air-pressure manipulation.”

4.2 — The output rehabilitates the method. The strongest argument for the protocol is not an argument; it is an undeniable artifact the protocol produced. A skeptic who can wave off a framework document cannot wave off a hand-proofed deliverable with reviewers and a filing date. Sequence accordingly: artifact first, method second. The deniable method is vouched for by the undeniable output.

4.3 — Contribution before initiation. A collaborator earns the deeper (method) layer by working the urgent (object) layer first — and the object work is itself the best onboarding to the method, because the collaborator experiences the output quality before being told how it was produced.

4.4 — The convergence is genealogy, not magic. When the machine collaborator’s instincts match the winning real-world standards, that is common cause, not coincidence: legibility drives adoption drives documentation drives training density drives fluency. The operator’s check on this is the bias it implies — the machine over-reaches for the popular standard even where a quieter, better one exists; domain-tacit knowledge (the standards professionals know but the internet under-documents) is the operator’s to supply.

Part Five — Retired Riders (documented, not core)

Per the armor-kit principle — *adopt the pattern, drop the vestige* — the following early scaffolding is retired from active protocol status and preserved here as documented alternatives for operators whose contexts genuinely need them:

- **Heavy per-exchange ritual.** Early framings implied fixed ceremonial structure around exchanges. The field run showed the protocol’s value lives in the *canon, the sweep, and the dialect*, not in ritual; session mechanics should stay as light as the work allows. Retained as an option for multi-operator or formal-records contexts.

- **Single-substrate assumptions.** Any language tying the method to one platform’s memory, context, or continuity features is retired; the four-artifact transfer supersedes it. The protocol is substrate-independent by design.
- **Validation-by-collaborator.** Any implication that the resident AI’s agreement validates the work is retired and reversed: the resident collaborator is explicitly *not* the check on its own output. Validation is external and multi-sourced (Part Two, 2.3).

Nothing retired is deleted; a future steward with a different threat model or team structure may re-adopt any of it deliberately.

Closing note

Version 2 asked whether dense understanding could move between a mind and a record at thought-speed. Version 3 answers with the extensions above — field-derived from a multi-week production run under load, and formalized as protocol only after the run completed. The core did not need to change. It needed to be run — and running it is what produced everything above. The method is not finished, and is not expected to be; the protocol now includes the mechanism for carrying itself forward.

Core frozen. Extensions accreted. Field-tested. — v3.1

Revision v3.1 (July 2026): field-report language revised to stand independent of adjacent projects. No core changes.

End of protocol extension.

Andrew Morton

July 2026 — Palmer Lake, Colorado